



PARALLEL COMPUTING WITH DASK

# Understanding Computer Storage & Big Data

**Dhavide Aruliah**

Director of Training, Anaconda

# "Data > one machine"



# Storage Units: Bytes, Kilobytes, Megabytes, ...

watt	W	
Kilowatt	KW	$10^3$ W
Megawatt	MW	$10^6$ W
Gigawatt	GW	$10^9$ W
Terawatt	TW	$10^{12}$ W

- Conventional units: factors of 1000
  - Kilo  $\rightarrow$  Mega  $\rightarrow$  Giga  $\rightarrow$  Tera  $\rightarrow$  ...

Byte	B	$2^3$ bits
Kilobyte	KB	$2^{10}$ Bytes
Megabyte	MB	$2^{20}$ Bytes
Gigabyte	GB	$2^{30}$ Bytes
Terabyte	TB	$2^{40}$ Bytes

- Binary computers: base 2:
  - Binary digit (bit)
  - Byte:  $2^3$  bits = 8 bits
  - $10^3 = 1000 \mapsto 2^{10} = 1024$




# Hard disks



- Hard storage: hard disks (permanent, big, **slow**)



# Random Access Memory (RAM)



- Soft storage: RAM (temporary, small, **fast**)





# Time Scales of Storage Technologies


Storage medium	Access time
RAM	120 ns
Solid-state disk	50-150 $\mu$ s
Rotational disk	1-10 ms
Internet (SF to NY)	40 ms

Storage medium	Rescaled
RAM	1 s
Solid-state disk	7-21 min
Rotational disk	2.5 hr - 1 day
Internet (SF to NY)	3.9 days



# Big Data in Practical Terms

- RAM: **fast** (ns- $\mu$ s)
- Hard disk: **slow** ( $\mu$ s-ms)
- I/O (input/output) is **punitive!**



# Querying Python interpreter's Memory Usage

```
In [1]: import psutil, os
```

```
In [2]: def memory_footprint():  
...:     '''Returns memory (in MB) being used by Python process'''  
...:     mem = psutil.Process(os.getpid()).memory_info().rss  
...:     return (mem / 1024 ** 2)
```



# Allocating Memory for an Array

```
In [3]: import numpy as np

In [4]: before = memory_footprint()

In [5]: N = (1024 ** 2) // 8 # Number of floats that fill 1 MB

In [6]: x = np.random.randn(50*N) # Random array filling 50 MB

In [7]: after = memory_footprint()

In [8]: print('Memory before: {} MB'.format(before))
Memory before: 45.68359375 MB

In [9]: print('Memory after: {} MB'.format(after))
Memory after: 95.765625 MB
```

# Allocating Memory for a Computation

```
In [10]: before = memory_footprint()
```

```
In [11]: x ** 2 # Computes, but doesn't bind result to a variable
```

```
Out[11]:
```

```
array([ 0.16344891,  0.05993282,  0.53595334, ...,  0.50537523,  
        0.48967157,  0.06905984])
```

```
In [12]: after = memory_footprint()
```

```
In [13]: print('Extra memory obtained: {} MB'.format(after - before))
```

```
Extra memory obtained: 50.34375 MB
```



# Querying Array Memory Usage

```
In [14]: x.nbytes # memory footprint in bytes (B)
```

```
Out[14]: 52428800
```

```
In [15]: x.nbytes // (1024**2) # memory footprint in megabytes (MB)
```

```
Out[15]: 50
```

# Querying DataFrame Memory Usage

```
In [16]: df = pd.DataFrame(x)
```

```
In [17]: df.memory_usage(index=False)
```

```
Out[17]:
```

```
0      52428800
```

```
dtype: int64
```

```
In [18]: df.memory_usage(index=False) // (1024**2)
```

```
Out[18]:
```

```
0         50
```

```
dtype: int64
```



## PARALLEL COMPUTING WITH DASK

**Let's practice!**



PARALLEL COMPUTING WITH DASK

# Thinking about Data in Chunks

**Dhavide Aruliah**

Director of Training, Anaconda

# Using `pd.read_csv` with `chunksize`

```
In [3]: filename = 'NYC_taxi_2013_01.csv'

In [4]: for chunk in pd.read_csv(filename, chunksize=50000):
...:     print('type: %s shape %s' %
...:           (type(chunk), chunk.shape))
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (50000, 14)
type: <class 'pandas.core.frame.DataFrame'> shape (49999, 14)
```



# Examining a chunk

```
In [5]: chunk.shape
```

```
Out[5]: (49999, 14)
```

```
In [6]: chunk.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 49999 entries, 150000 to 199998
```

```
Data columns (total 14 columns):
```

medallion	49999	non-null	object
hack_license	49999	non-null	object
vendor_id	49999	non-null	object
rate_code	49999	non-null	int64
store_and_fwd_flag	162	non-null	object
pickup_datetime	49999	non-null	object
dropoff_datetime	49999	non-null	object
passenger_count	49999	non-null	int64
trip_time_in_secs	49999	non-null	int64
trip_distance	49999	non-null	float64
pickup_longitude	49999	non-null	float64
pickup_latitude	49999	non-null	float64
dropoff_longitude	49999	non-null	float64
dropoff_latitude	49999	non-null	float64

```
dtypes: float64(5), int64(3), object(6)  
memory usage: 5.3+ MB
```

# Filtering a chunk

```
In [7]: is_long_trip = (chunk.trip_time_in_secs > 1200)
```

```
In [8]: chunk.loc[is_long_trip].shape
```

```
Out[8]: (5565, 14)
```

	passenger_count	trip_time_in_secs	trip_distance
167	1	300	2.1
168	3	2100	13.51
169	1	420	1.56
170	3	120	0.67
171	4	960	3.34
172	2	1140	4.13
173	5	300	2.19
174	1	1620	10.1
175	1	120	0.55
176	1	1440	10.63
177	1	120	0.47
178	1	1320	6.82
179	1	1500	5.32
180	1	420	1.71
181	3	960	4.72
182	6	1020	4.77
183	1	600	1.73
184	1	1020	7.29
185	3	1260	11.17

	passenger_count	trip_time_in_secs	trip_distance
168	3	2100	13.51
174	1	1620	10.1
176	1	1440	10.63
178	1	1320	6.82
179	1	1500	5.32
185	3	1260	11.17

# Chunking & filtering together

```
In [9]: def filter_is_long_trip(data):  
...:     "Returns DataFrame filtering trips longer than 20 minutes"  
...:     is_long_trip = (data.trip_time_in_secs > 1200)  
...:     return data.loc[is_long_trip]
```

```
In [10]: chunks = []
```

```
In [11]: for chunk in pd.read_csv(filename, chunksize=1000):  
...:     chunks.append(filter_is_long_trip(chunk))
```

```
In [12]: chunks = [filter_is_long_trip(chunk)  
...:                 for chunk in pd.read_csv(filename,  
...:                 chunksize=1000) ]
```

# Using `pd.concat()`

```
In [13]: len(chunks)
```

```
Out[13]: 200
```

```
In [14]: lengths = [len(chunk) for chunk in chunks]
```


```
In [15]: lengths[-5:] # each has ~100 rows
```

```
Out[15]: [115, 147, 137, 109, 119]
```

```
In [16]: long_trips_df = pd.concat(chunks)
```

```
In [17]: long_trips_df.shape
```

```
Out[17]: (21661, 14)
```





# Plotting the filtered results

```
In [18]: import matplotlib.pyplot as plt

In [19]: long_trips_df.plot.scatter(x='trip_time_in_secs',
    ...:                             y='trip_distance');

In [20]: plt.xlabel('Trip duration [seconds]');

In [21]: plt.ylabel('Trip distance [miles]');

In [22]: plt.title('NYC Taxi rides over 20 minutes (2013-01-01 to 2013-01-14)');

In [23]: plt.show();
```



## PARALLEL COMPUTING WITH DASK

**Let's practice!**





PARALLEL COMPUTING WITH DASK

# Managing Data with Generators

Dhavide Aruliah

Director of Training, Anaconda

# Filtering in a List Comprehension

```
In [1]: import pandas as pd

In [2]: filename = 'NYC_taxi_2013_01.csv'

In [3]: def filter_is_long_trip(data):
...:     "Returns DataFrame filtering trips longer than 20 minutes"
...:     is_long_trip = (data.trip_time_in_secs > 1200)
...:     return data.loc[is_long_trip]

In [4]: chunks = [filter_is_long_trip(chunk)
...:                 for chunk in pd.read_csv(filename, chunksize=1000)]
```

# Filtering & Summing with Generators

```
In [5]: chunks = (filter_is_long_trip(chunk)
...:               for chunk in pd.read_csv(filename, chunksize=1000))

In [6]: distances = (chunk['trip_distance'].sum() for chunk in chunks)

In [7]: sum(distances)
Out[7]: 230909.56000000003
```



# Examining Consumed Generators

```
In [8]: distances
```

```
Out[8]: <generator object <genexpr> at 0x10766f9e8>
```

```
In [9]: next(distances)
```

```
StopIteration
```

```
Traceback (most recent call last)
```

```
<ipython-input-10-9995a5373b05> in <module>()
```

# Reading Many Files

```
In [10]: template = 'yellow_tripdata_2015-{:02d}.csv'

In [11]: filenames = (template.format(k) for k in range(1,13)) # generator

In [12]: for fname in filenames:
...:     print(fname) # Examine contents
yellow_tripdata_2015-01.csv
yellow_tripdata_2015-02.csv
yellow_tripdata_2015-03.csv
yellow_tripdata_2015-04.csv
yellow_tripdata_2015-05.csv
yellow_tripdata_2015-06.csv
yellow_tripdata_2015-07.csv
yellow_tripdata_2015-08.csv
yellow_tripdata_2015-09.csv
yellow_tripdata_2015-10.csv
yellow_tripdata_2015-11.csv
yellow_tripdata_2015-12.csv
```

# Examining a Sample DataFrame

```
In [16]: df = pd.read_csv('yellow_tripdata_2015-12.csv', parse_dates=[1, 2])
```

```
In [17]: df.info() # columns deleted from output
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 71634 entries, 0 to 71633
```

```
Data columns (total 19 columns):
```

```
VendorID          71634 non-null int64
```

```
tpep_pickup_datetime  71634 non-null datetime64[ns]
```

```
tpep_dropoff_datetime 71634 non-null datetime64[ns]
```

```
passenger_count      71634 non-null int64
```

```
...
```

```
...
```

```
dtypes: datetime64[ns](2), float64(12), int64(4), object(1)
```

```
memory usage: 10.4+ MB
```

```
In [18]: def count_long_trips(df):
```

```
    ....:     df['duration'] = (df.tpep_dropoff_datetime -
    ....:                        df.tpep_pickup_datetime).dt.seconds
```

```
    ....:     is_long_trip = df.duration > 1200
```

```
    ....:     result_dict = {'n_long': [sum(is_long_trip)],
    ....:                    'n_total': [len(df)]}
```

```
    ....:     return pd.DataFrame(result_dict)
```

# Aggregating with Generators

```
In [18]: def count_long_trips(df):
...:     df['duration'] = (df.tpep_dropoff_datetime -
...:                      df.tpep_pickup_datetime).dt.seconds
...:     is_long_trip = df.duration > 1200
...:     result_dict = {'n_long': [sum(is_long_trip)],
...:                    'n_total': [len(df)]}
...:     return pd.DataFrame(result_dict)

In [19]: filenames = [template.format(k) for k in range(1,13)] # listcomp

In [20]: dataframes = (pd.read_csv(fname, parse_dates=[1,2])
...:                    for fname in filenames) # generator

In [21]: totals = (count_long_trips(df) for df in dataframes) # generator

In [22]: annual_totals = sum(totals) # Consumes generators
```



# Computing the Fraction of Long Trips

```
In [23]: print(annual_totals)
         n_long  n_total
0  172617    851390

In [24]: fraction = annual_totals['n_long'] / annual_totals['n_total']

In [25]: print(fraction)
0      0.202747
dtype: float64
```



## PARALLEL COMPUTING WITH DASK

**Let's practice!**



PARALLEL COMPUTING WITH DASK

# Delaying Computation with Dask

**Dhavide Aruliah**

Director of Training, Anaconda

# Composing functions

```
In [1]: from math import sqrt
```

```
In [2]: def f(z):  
...:     return sqrt(z + 4)
```

```
In [3]: def g(y):  
...:     return y - 3
```

```
In [4]: def h(x):  
...:     return x ** 2
```

```
In [5]: x = 4
```

```
In [6]: y = h(x)
```

```
In [7]: z = g(y)
```

```
In [8]: w = f(z)
```

```
In [9]: print(w) # final result  
4.123105625617661
```

```
In [10]: print(f(g(h(x)))) # equivalent to before  
4.123105625617661
```

# Deferring Computation with delayed

```
In [11]: from dask import delayed

In [12]: y = delayed(h)(x)

In [13]: z = delayed(g)(y)

In [14]: w = delayed(f)(z)

In [15]: print(w)
Delayed('f-5f9307e5-eb43-4304-877f-1df5c583c11c')


In [16]: type(w) # a dask Delayed object
Out[16]: dask.delayed.Delayed

In [17]: w.compute() # computation occurs *now*
Out[17]: 4.123105625617661
```



# Visualizing a Task Graph

```
In [18]: w.visualize()
```



# Renaming Decorated Functions

```
In [19]: f = delayed(f)

In [20]: g = delayed(g)

In [21]: h = delayed(h)

In [22]: w = f(g(h(4)))

In [23]: type(w) # a dask Delayed object
Out[23]: dask.delayed.Delayed

In [24]: w.compute() # computation occurs *now*
Out[24]: 4.123105625617661
```



# Using Decorator @-Notation

```
In [25]: def f(x):  
...:     return sqrt(x + 4)  
  
In [26]: f = delayed(f)  
  
In [27]: @delayed # Equivalent to definition in above 2 cells  
...: def f(x):  
...:     return sqrt(x + 4)
```

# Deferring Computation with Loops

```
In [28]: @delayed
...: def increment(x):
...:     return x + 1
```

```
In [29]: @delayed
...: def double(x):
...:     return 2 * x
```

```
In [30]: @delayed
...: def add(x, y):
...:     return x + y
```

```
In [31]: data = [1, 2, 3, 4, 5]
```

```
In [32]: output = []
```

```
In [33]: for x in data:
...:     a = increment(x)
...:     b = double(x)
...:     c = add(a, b)
...:     output.append(c)
```

```
In [34]: total = sum(output)
```


```
In [35]: total
Out[35]: Delayed('add-c6803f9e890c95cec8e2e3dd3c62b384')
```

```
In [36]: output
Out[36]:
[Delayed('add-6a624d8b-8ddb-44fc-b0f0-0957064f54b7'),
 Delayed('add-9e779958-f3a0-48c7-a558-ce47fc9899f6'),
 Delayed('add-f3552c6f-b09d-4679-a770-a7372e2c278b'),
 Delayed('add-ce05d7e9-42ec-4249-9fd3-61989d9a9f7d'),
 Delayed('add-dd950ec2-c17d-4e62-a267-1dabe2101bc4')]
```

```
In [37]: total.visualize()
```



# Visualizing the task graph



# Aggregating with delayed Functions

```
In [35]: template = 'yellow_tripdata_2015-{:02d}.csv'

In [36]: filenames = [template.format(k) for k in range(1,13)]

In [37]: @delayed
...: def count_long_trips(df):
...:     df['duration'] = (df.tpep_dropoff_datetime -
...:                      df.tpep_pickup_datetime).dt.seconds
...:     is_long_trip = df.duration > 1200
...:     result_dict = {'n_long': [sum(is_long_trip)],
...:                   'n_total': [len(df)]}
...:     return pd.DataFrame(result_dict)

In [38]: @delayed
...: def read_file(fname):
...:     return pd.read_csv(fname, parse_dates=[1,2])
```

# Computing Fraction of Long Trips with delayed Functions

```
In [39]: totals = [count_long_trips(read_file(fname))  
....:                                     for fname in filenames]
```

```
In [40]: annual_totals = sum(totals)
```

```
In [41]: annual_totals = annual_totals.compute()
```

```
Out[41]:  
   n_long  n_total  
0  172617   851390
```

```
In [42]: fraction = annual_totals['n_long'] /  
....: annual_totals['n_total']
```

```
In [43]: print(fraction)  
0      0.202747  
dtype: float64
```



## PARALLEL COMPUTING WITH DASK

**Let's practice!**